

LAB: String Patterns, Sorting & Grouping

The practice problems for this Lab will provide hands on experience with string patterns, sorting result sets and grouping result sets. You will also learn how to run SQL scripts to create several tables at once, as well as how to load data into tables from .csv files.

HR Database

We will be working on a sample HR database for this Lab. This HR database schema consists of 5 tables called EMPLOYEES, JOB_HISTORY, JOBS, DEPARTMENTS and LOCATIONS. Each table has a few rows of sample data. The following diagram shows the tables for the HR database.

SAMPLE HR DATABASE TABLES

EMPLOYEES

EMP_ID	F_NAME	L_NAME	SSN	B_DATE	SEX	ADDRESS	JOB_ID	SALARY	MANAGER_ID	DEP_ID
E1001	John	Thomas	123456	1976-01-09	M	5631 Rice, OakPark,IL	100	100000	30001	2
E1002	Alice	James	123457	1972-07-31	F	980 Berry Ln, Elgin,IL	200	80000	30002	5
E1003	Steve	Wells	123458	1980-08-10	M	291 Springs, Gary,IL	300	50000	30002	5

JOB_HISTORY

EMPL_ID	START_DATE	JOBS_ID	DEPT_ID
E1001	2000-01-30	100	2
E1002	2010-08-16	200	5
E1003	2016-08-10	300	5

JOBS

JOB_IDENT	JOB_TITLE	MIN_SALARY	MAX_SALARY
100	Sr. Architect	60000	100000
200	Sr.SoftwareDeveloper	60000	80000
300	Jr.SoftwareDeveloper	40000	60000

DEPARTMENTS

DEPT_ID_DEP	DEP_NAME	MANAGER_ID	LOC_ID
2	Architect Group	30001	L0001
5	Software Development	30002	L0002
7	Design Team	30003	L0003
5	Software	30004	L0004

LOCATIONS

LOCT_ID	DEP_ID_LOC
L0001	2
L0002	5
L0003	7

To complete this lab you will utilize Db2 database service on IBM Cloud as you did for the previous lab. There are three parts to this lab:

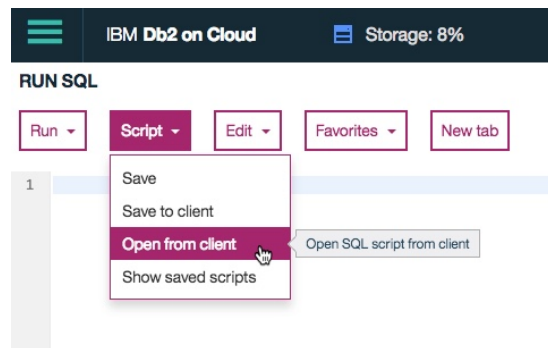
- I. Creating tables
- II. Loading data into tables
- III. Composing and running queries

Part I: CREATING TABLES

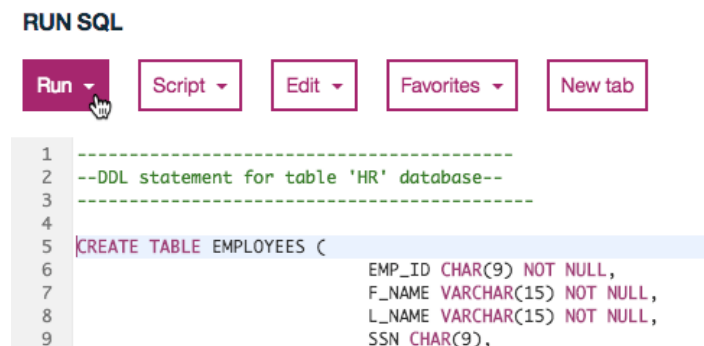
If you do not yet have access to Db2 on IBM Cloud, please refer to Lab Instructions in the Module/Week 1.

Rather than create each table manually by typing the DDL commands in the SQL editor, you will execute a script containing the create table commands for all the tables. Following step by step instructions are provided to perform this:

- 1) Download the script file "Script_Create_Tables.sql" provided on the Lab page
- 2) Login to IBM Cloud and go to the Dashboard:
<https://console.bluemix.net/dashboard/apps> where you can find the Db2 service that you created in a previous Lab. Click on the service. Next, open the Db2 Console by clicking on "Open Console" button. Go to the Run SQL page. The Run SQL tool enables you to run DDL and SQL statements.
- 3) Click on the Folder icon to "Open script". Locate the file Script_Create_Tables.sql that you downloaded to your computer earlier and open it.



- 4) Once the statements are in the SQL Editor tool, you can run the queries against the database by selecting the option "Run" then "Run All".



5) At the bottom of the screen you will see a Result section. Clicking on a query will show the execution details of the job - whether it ran successfully, or had any errors or warnings. Ensure your queries ran successfully and created all the tables.

Saved scripts **Result**

Filter by status: All Delete All

▼ All(5), Failed(0)

- ✓ -----DDL statement for table '...
- ✓ CREATE TABLE JOB_HISTORY (EMPL_ID CHAR(9) NOT NULL, ...
- ✓ CREATE TABLE JOBS (JOB_IDENT CHAR(9) NOT NULL, JOB_T...
- ✓ CREATE TABLE DEPARTMENTS (DEPT_ID_DEP CHAR(9) NOT ...
- ✓ CREATE TABLE LOCATIONS (LOCT_ID CHAR(9) NOT NULL, D...

Log

The statement ran successfully.

SQL statement

```
CREATE TABLE LOCATIONS (
  LOCT_ID CHAR(9) NOT NULL,
  DEP_ID_LOC CHAR(9) NOT NULL,
  PRIMARY KEY (LOCT_ID, DEP_ID_L)
```


Execution log

Run time: 0.103 s
Status: SUCCEEDED

6) You can delete the Jobs history by clicking on the Delete icon as shown below.

Jobs Clear All

Finished successfully ▼

All (5) | Failed (0) | 

-----DDL statement for table 'HR' datab...

```
CREATE TABLE JOB_HISTORY ( EMPL_ID CHAR(9) NOT NULL, START_DA...
CREATE TABLE JOBS ( JOB_IDENT CHAR(9) NOT NULL, JOB_TITLE VARC...
CREATE TABLE DEPARTMENTS ( DEPT_ID_DEP CHAR(9) NOT NULL, DEP_...
CREATE TABLE LOCATIONS ( LOCT_ID CHAR(9) NOT NULL, DEP_ID_LOC ...
```

7) Now you can look at the tables you created. Navigate to the three bar menu icon and select "Explore". Select the correct Schema to see the newly created tables.

EXPLORE

IBM Db2 on Cloud Storage: 27% Discover Refresh

☐ Show system schemas

Schema	Table	Table definition															
QCM54853	DEPARTMENTS	<p>DEPARTMENTS</p> <p>No statistics available.</p> <table border="1"> <thead> <tr> <th>COLUMN NAME</th> <th>DATA TYPE</th> <th>NULLABLE</th> </tr> </thead> <tbody> <tr> <td>DEPT_ID_...</td> <td>CHARA...</td> <td>N</td> </tr> <tr> <td>DEP_NAME</td> <td>VARCHAR</td> <td>Y</td> </tr> <tr> <td>MANAGE...</td> <td>CHARA...</td> <td>Y</td> </tr> <tr> <td>LOC_ID</td> <td>CHARA...</td> <td>Y</td> </tr> </tbody> </table>	COLUMN NAME	DATA TYPE	NULLABLE	DEPT_ID_...	CHARA...	N	DEP_NAME	VARCHAR	Y	MANAGE...	CHARA...	Y	LOC_ID	CHARA...	Y
COLUMN NAME	DATA TYPE	NULLABLE															
DEPT_ID_...	CHARA...	N															
DEP_NAME	VARCHAR	Y															
MANAGE...	CHARA...	Y															
LOC_ID	CHARA...	Y															
ERRORSCHEMA Sample	EMPLOYEES																
ST_INFORMTN_SCHEMA Sample	INSTRUCTOR																
	JOBS																
	JOB_HISTORY																
	LOCATIONS																

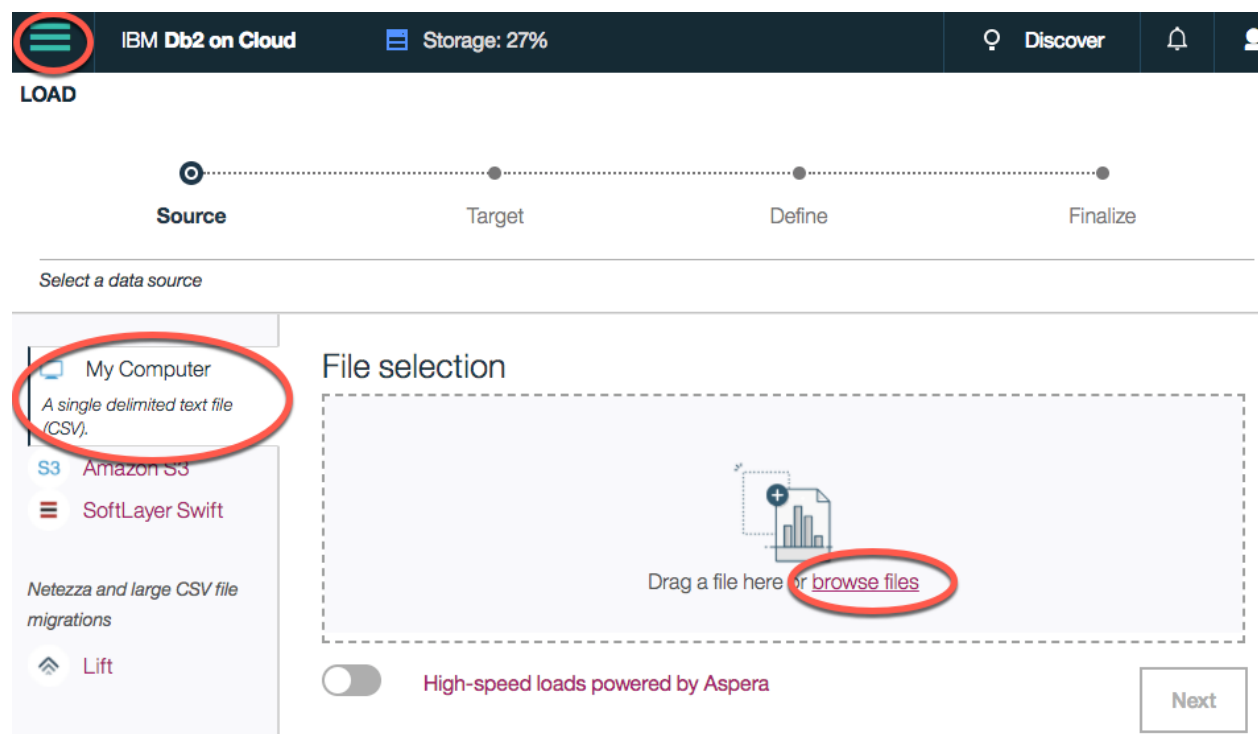
Part II: LOADING DATA

Now let us see how data can be loaded into Db2. We could manually insert each row into the table one by one but that would take a long time. Instead, Db2 (and almost every other database) allows you to Load data from csv files.

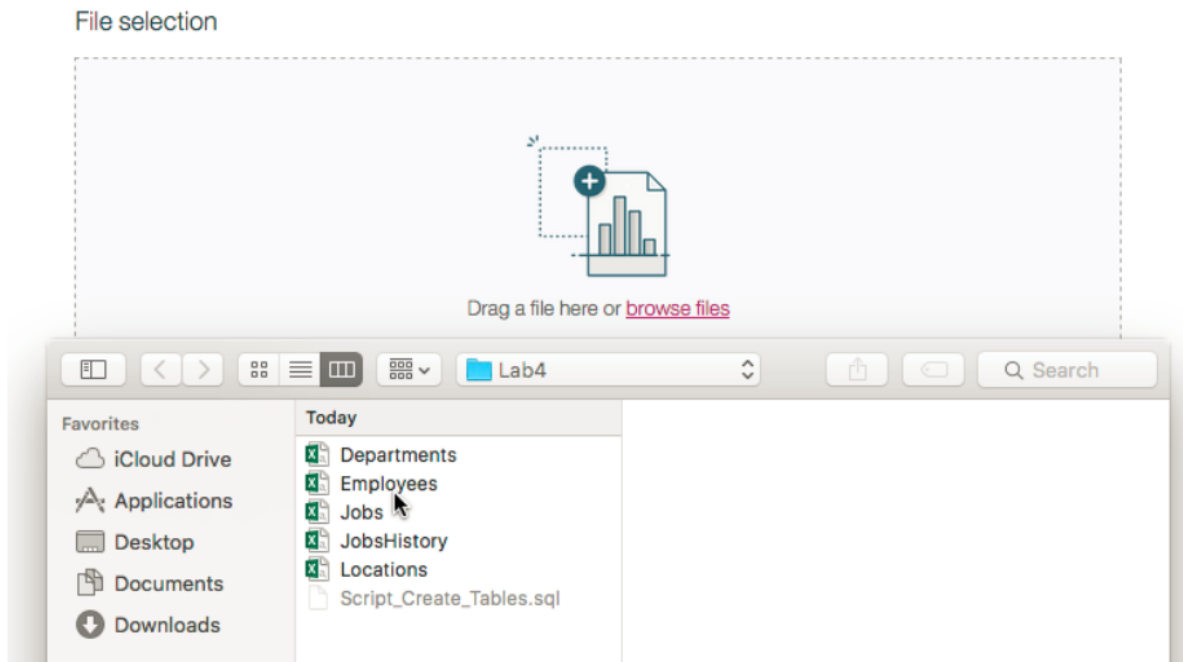
Please follow the steps below which explains the process of loading data into the tables we created earlier.

1) Download the 5 required data source files:
("Employees.csv", "Departments.csv", "Jobs.csv", "JobsHistory.csv", "Locations.csv") to your computer:

2) First let us learn how to load data into the Employees table that we created earlier. From the 3 bar menu icon, navigate to the "Load" page and ensure "My Computer" is selected as the source. Click on the "browse files" link.

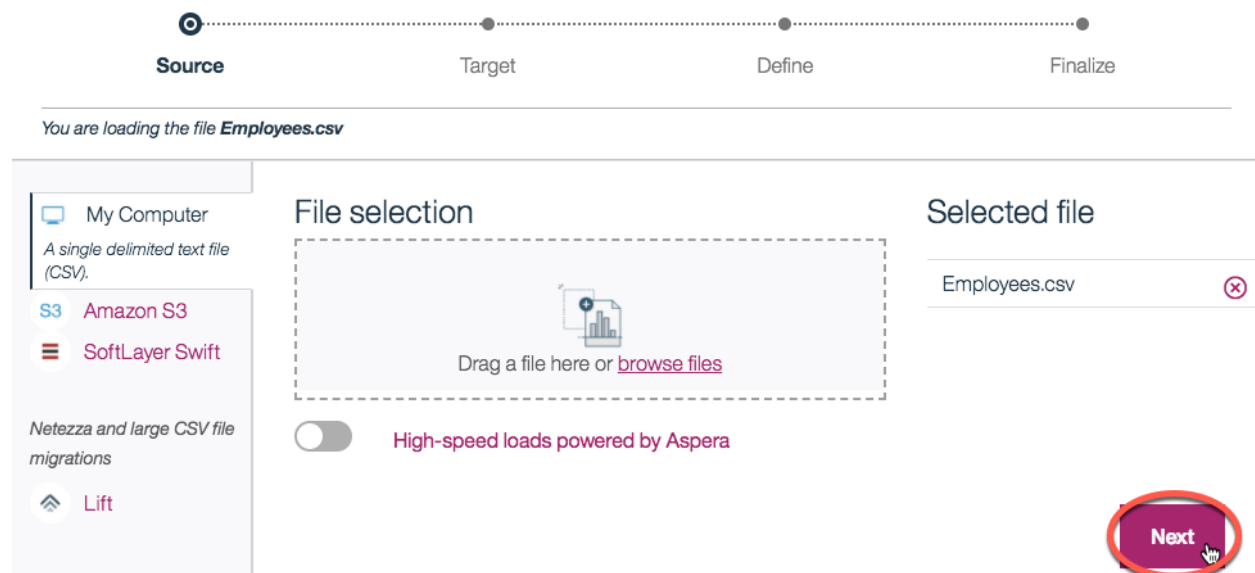


3) Choose the file "Employees.csv" that you downloaded to your computer and click "Open".



~~4) Once the File is selected click "Next" in the bottom right corner.~~

LOAD



~~5) Select the schema for your Userid. It will show all the tables that have been created in this Schema previously, including the Employees table. Select the EMPLOYEES table.~~

You are loading the file **Employees.csv** into **QCM54853.EMPLOYEES**

Select a load target

Schema	Table	Table definition									
Find a schema	Find a table in QCM54853	EMPLOYEES Approximate 0 rows (0 KB) Updated on 5/2/2018 at 1:09:31 PM									
QCM54853	DEPARTMENTS	<input checked="" type="radio"/> Append new data <input type="radio"/> Overwrite table with new data									
ERRORSCHEMA <i>Sample</i>	EMPLOYEES	<table border="1"> <thead> <tr> <th>COLUMN NAME</th> <th>DATA TYPE</th> <th>NULLABLE</th> </tr> </thead> <tbody> <tr> <td>EMP_ID</td> <td>CHARA...</td> <td></td> </tr> <tr> <td>F_NAME</td> <td>VARCHAR</td> <td></td> </tr> </tbody> </table>	COLUMN NAME	DATA TYPE	NULLABLE	EMP_ID	CHARA...		F_NAME	VARCHAR	
COLUMN NAME	DATA TYPE	NULLABLE									
EMP_ID	CHARA...										
F_NAME	VARCHAR										
ST_INFORMTN_SCHEMA <i>Sample</i>	INSTRUCTOR										
	JOBS										
	JOB_HISTORY										

[Back](#)
[Next](#)

6) Since our source data files do not contain any rows with column labels, turn off the setting for "Header in first row". Also, click on the down arrow next to "Date format" and choose "MM/DD/YYYY" since that is how the date is formatted in our source file.

You are loading the file **Employees.csv** into **QCM54853.EMPLOYEES**

Code page (character encoding): 1208 (UTF-8) Separator: , Header in first row: ☐ Time & date format: ?

Date format: MM/DD/YYYY Time format: HH:MM:SS Timestamp format: YYYY-MM-DD HH:MM:SS

	EMP_ID CHARACTER	F_NAME VARCHAR	L_NAME VARCHAR	SSN CHARACTER	B_DATE DATE	SEX CHARACTER	ADDRESS VARCHAR
1	E1001	John	Thomas	123456	01/09/1976	M	"5631 Rice
2	E1002	Alice	James	123457	07/31/1972	F	980 Berry Ln, Elgin,IL
3	E1003	Steve	Wells	123458	08/10/1980	M	291 Springs, Gary,IL
4	E1004	Santosh	Kumar	123459	07/20/1985	M	511 Aurora Av, Aurora,IL
5	E1005	Ahmed	Hussain	123410	01/04/1981	M	216 Oak Tree, Geneva,IL
6	E1006	Nancy	Allen	123411	02/06/1978	F	111 Green Pl, Elgin,IL
7	E1007	Mary	Thomas	123412	05/05/1975	F	100 Rose Pl, Gary,IL
8	E1008	Bharath	Gupta	123413	05/06/1985	M	145 Berry Ln, Naperville,IL
9	E1009	Andrea	Jones	123414	07/09/1990	F	120 Fall Creek, Gary,IL
10		Ann		123415	03/01/1982		10000 Berry Sp, Elgin,IL

[Back](#) [Next](#)

7) Click "Next". Review the Load setting and click "Begin Load" in the top-right corner.

Review settings

Summary

Code page: 1208 (Default)

Separator: , (Default)

Header in first row: No

Time format: HH:MM:SS (Default)

Date format: MM/DD/YYYY

Timestamp format: YYYY-MM-DD
HH:MM:SS (Default)

String delimiter: "(Default)

Option


Maximum number of warnings

1000

[Back](#)
[Begin Load](#)

8) After Loading is complete you will notice that we were successful in loading all 10 rows of the Employees table. If there are any Errors or Warnings you can view them on this screen.


Load details



My computer **Target**
Employees.csv QCM54853.EMPLOYEES

[View Table](#)
[Load More Data](#)

Status Settings



10 10 0

Rows read Rows loaded Rows rejected

Start time
05/02/2018 1:51:27 PM

End time
05/02/2018 1:51:28 PM

The data load job succeeded.

You can now work with your data.

Errors 0 **Warnings** 0

No errors

9) You can see the data that was loaded by clicking on the View Table. Alternatively you can go into the Explore page and page select the correct schema, then the EMPLOYEES table, and click "View Data".

QCM54853.EMPLOYEES

[Delete Table](#)
[Export to CSV](#)

	EMP_ID CHARACTER(9)	F_NAME VARCHAR(15)	L_NAME VARCHAR(15)	SSN CHARACTER(9)	B_DATE DATE	SEX CHARACTER(1)	ADDRESS VARCHAR(30)	JOB_ID CHARACTER(9)
1	E1001	John	Thomas	123456	1976-01-09	M	5631 Rice, OakPark,	100
2	E1002	Alice	James	123457	1972-07-31	F	980 Berry In, Elgin,IL	200
3	E1003	Steve	Wells	123458	1980-08-10	M	291 Springs, Gary,IL	300
4	E1004	Santosh	Kumar	123459	1985-07-20	M	511 Aurora Av, Auror	400
5	E1005	Ahmed	Hussain	123410	1981-01-04	M	216 Oak Tree, Gene	500
6	E1006	Nancy	Allen	123411	1978-02-06	F	111 Green Pl, Elgin,IL	600
7	E1007	Mary	Thomas	123412	1975-05-05	F	100 Rose Pl, Gary,IL	650
8	E1008	Bharath	Gupta	123413	1985-05-06	M	145 Berry Ln, Naper	660
9	E1009	Andrea	Jones	123414	1990-07-09	F	120 Fail Creek, Gary,	234
10	E1010	Ann	Jacob	123415	1982-03-30	F	111 Britany Springs,E	220

10. Now its your turn to load the remaining 4 tables of the HR database – Locations, JobHistory, Jobs, and Departments. Please follow the steps above to load the data from the remaining source files.

Question 1: Were there any warnings loading data into the JOBS table? What can be done to resolve this?



Hint: View the data loaded into this table and pay close attention to the JOB_TITLE column.



Question 2: Did all rows from the source file load successfully in the DEPARTMENT table? If not, are you able to figure out why not?



Hint: Look at the warning. Also, note the Primary Key for this table.



Part III: COMPOSING AND RUNNING QUERIES

You created the tables for the HR database schema and also learned how to load data into these tables. Now try and work on a few advanced DML queries that were introduced in this module.

Follow these steps to create and run the queries indicated below

- 1) Navigate to the Run SQL tool in Db2 on Cloud
- 2) Compose query and run it.
- 3) Check the Logs created under the Results section. Looking at the contents of the Log explains whether the SQL statement ran successfully. Also look at the query results to ensure the output is what you expected.

Query 1: Retrieve all employees whose address is in Elgin, IL

[Hint: Use the LIKE operator to find similar strings]

Query 2: Retrieve all employees who were born during the 1970's.

[Hint: Use the LIKE operator to find similar strings]

Query 3: Retrieve all employees in department 5 whose salary is between 60000 and 70000.

[Hint: Use the keyword BETWEEN for this query]

Query 4A: Retrieve a list of employees ordered by department ID.

[Hint: Use the ORDER BY clause for this query]

Query 4B: Retrieve a list of employees ordered in descending order by department ID and within each department ordered alphabetically in descending order by last name.

Query 5A: For each department ID retrieve the number of employees in the department.

[Hint: Use COUNT(*) to retrieve the total count of a column, and then GROUP BY]

Query 5B: For each department retrieve the number of employees in the department, and the average employees salary in the department.

[Hint: Use COUNT(*) to retrieve the total count of a column, and AVG() function to compute average salaries, and then group]

Query 5C: Label the computed columns in the result set of Query 5B as "NUM_EMPLOYEES" and "AVG_SALARY".

[Hint: Use AS "LABEL_NAME" after the column name]

Query 5D: In Query 5C order the result set by Average Salary.

[Hint: Use ORDER BY after the GROUP BY]

Query 5E: In Query 5D limit the result to departments with fewer than 4 employees.

[Hint: Use HAVING after the GROUP BY, and use the count() function in the HAVING clause instead of the column label.

Note: WHERE clause is used for filtering the entire result set whereas the HAVING clause is used for filtering the result of the grouping]

BONUS Query 6: Similar to 4B but instead of department ID use department name. Retrieve a list of employees ordered by department name, and within each department ordered alphabetically in descending order by last name.

[Hint: Department name is in the DEPARTMENTS table. So your query will need to retrieve data from more than one table. Don't worry if you are not able to figure this one out ... we'll cover working with multiple tables in the next lesson.]

In this lab you learned how to work with string patterns, sorting result sets and grouping result sets.

Thank you for completing this lab!

See solutions on the following page.

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

```

1  -- Query 1-----
2  ;
3  select F_NAME , L_NAME
4  from EMPLOYEES
5  where ADDRESS LIKE '%Elgin,IL%';
6  --Query 2--
7  ;

```

Saved scripts **Result**

Filter by status: **Result set** Log

All

Delete All

▼ AI...

F_NAME	L_NAME
Alice	James
Nancy	Allen
Ann	Jacob

select...
select...
select...
Total rows: 3

[REDACTED]

RUN SQL

Run Script Edit Favorites New tab

```

6  --Query 2--
7  ;
8  select F_NAME , L_NAME
9  from EMPLOYEES
10 where B_DATE LIKE '197%';
11 ---Query3--
12 :

```

Saved scripts **Result**

Filter by status: **Result set** Log

All

Delete All

▼ AI...

F_NAME	L_NAME
John	Thomas
Alice	James
Nancy	Allen
Mary	Thomas

select...
select...
select...
select...
Total rows: 4

```

11 ---Query3---
12 ;
13 select *
14 from EMPLOYEES
15 where (SALARY BETWEEN 60000 and 70000) and DEP_ID = 5 ;
16 ---Query4---
17 ;

```

Saved scripts **Result**

Filter by status: **Result set** Log

All

Delete All

✓ All(5)...

✓ select F_...

✓ select F_...

✓ select * f...

EMP_ID	F_NAME	L_NAME	SSN	B_DATE	SEX	ADDRESS	JOB_ID	SALARY	MANAGER_ID	DEP_ID
E1004	Santosh	Kumar	1234...	1985-07-20	M	511 Aurora ...	400	60000.00	30004	5
E1010	Ann	Jacob	12341...	1982-03-30	F	111 Britany ...	220	70000.00	30004	5

Total rows: 2

```

select F_NAME, L_NAME, DEP_ID
from EMPLOYEES
order by DEP_ID;

```

red scripts **Result**

er by status: **Result set** Log

All

delete All

(1)...

select F_...

(1)...

select F_...

(1)...

select F_...

(1)...

select F_...

(1)...

select F_...

F_NAME	L_NAME	DEP_ID
John	Thomas	2
Ahmed	Hussain	2
Nancy	Allen	2
Alice	James	5
Steve	Wells	5
Santosh	Kumar	5
Ann	Jacob	5
Mary	Thomas	7
Bharath	Gupta	7
Andrea	Jones	7

Total rows: 10


```
1 select DEP_ID, COUNT(*), AVG(SALARY)
2 from EMPLOYEES
3 group by DEP_ID;
```

Saved scripts	Result
Script 1	Success
Script 2	Success
Script 3	Success
Script 4	Success
Script 5	Success
Script 6	Success
Script 7	Success
Script 8	Success
Script 9	Success
Script 10	Success
Script 11	Success
Script 12	Success
Script 13	Success
Script 14	Success
Script 15	Success
Script 16	Success
Script 17	Success
Script 18	Success
Script 19	Success
Script 20	Success
Script 21	Success
Script 22	Success
Script 23	Success
Script 24	Success
Script 25	Success
Script 26	Success
Script 27	Success
Script 28	Success
Script 29	Success
Script 30	Success
Script 31	Success
Script 32	Success
Script 33	Success
Script 34	Success
Script 35	Success
Script 36	Success
Script 37	Success
Script 38	Success
Script 39	Success
Script 40	Success
Script 41	Success
Script 42	Success
Script 43	Success
Script 44	Success
Script 45	Success
Script 46	Success
Script 47	Success
Script 48	Success
Script 49	Success
Script 50	Success
Script 51	Success
Script 52	Success
Script 53	Success
Script 54	Success
Script 55	Success
Script 56	Success
Script 57	Success
Script 58	Success
Script 59	Success
Script 60	Success
Script 61	Success
Script 62	Success
Script 63	Success
Script 64	Success
Script 65	Success
Script 66	Success
Script 67	Success
Script 68	Success
Script 69	Success
Script 70	Success
Script 71	Success
Script 72	Success
Script 73	Success
Script 74	Success
Script 75	Success
Script 76	Success
Script 77	Success
Script 78	Success
Script 79	Success
Script 80	Success
Script 81	Success
Script 82	Success
Script 83	Success
Script 84	Success
Script 85	Success
Script 86	Success
Script 87	Success
Script 88	Success
Script 89	Success
Script 90	Success
Script 91	Success
Script 92	Success
Script 93	Success
Script 94	Success
Script 95	Success
Script 96	Success
Script 97	Success
Script 98	Success
Script 99	Success
Script 100	Success

Filter by stat **Result set** LogAll

Delete All

AI

☒ sel.

DEP_ID	2	3
2	3	86666.6666666666666666666666666666
5	4	65000.0000000000000000000000000000
7	3	66666.6666666666666666666666666666

Total rows: 3

```
select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID;
```

ed scripts	Result
------------	--------

or by status: **Result set** Log

All 

ete All



select...



select ...

DEP_ID	NUM_EMPLOYEES	AVG_SALARY
2	3	86666.6666666666666666666666666666
5	4	65000.0000000000000000000000000000
7	3	66666.6666666666666666666666666666

Total rows: 3

```
select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID
order by AVG_SALARY;
```

red scripts	Result
-------------	--------

er by status: **Result set** Log

All

ete All



select...

select ...

DEP_ID	NUM_EMPLOYEES	AVG_SALARY
5	4	65000.000000000000000000000000
7	3	66666.666666666666666666666666
2	3	86666.666666666666666666666666

Total rows: 3

Query 5E: Output

```
select DEP_ID, COUNT(*) AS "NUM_EMPLOYEES", AVG(SALARY) AS "AVG_SALARY"
from EMPLOYEES
group by DEP_ID
having count(*) < 4
order by AVG_SALARY;
```

d scripts	Result
by status:	Result set Log
Delete All	
), F...	
select DEP_...	
select DEP_I...	
), F...	
	Total rows: 2

DEP_ID	NUM_EMPLOYEES
7	3
2	3

BONUS Query 6: Output

Note that in the Query below “D” and “E” are aliases for the table names. Once you define an alias like “D” in your query, you can simply write “D.COLUMN_NAME” rather than the full form ‘DEPARTMENTS.COLUMN_NAME’.

```
16 --Query4--
17 ;
18 select D.DEP_NAME , E.F_NAME, E.L_NAME
19 from EMPLOYEES as E, DEPARTMENTS as D
20 where E.DEP_ID = D.DEPT_ID_DEP
21 order by D.DEP_NAME, E.L_NAME desc ;
22 --Query5--
23 --
```

Saved scripts	Result
Filter by status:	Result set Log
All	
Delete All	
✓ All(5)...	
✓ select F_...	
✓ select F_...	
✓ select * fr...	
✓ select D...	
✓ select DE...	
	Total rows: 10

DEP_NAME	F_NAME	L_NAME
Architect Group	John	Thomas
Architect Group	Ahmed	Hussain
Architect Group	Nancy	Allen
Design Team	Mary	Thomas
Design Team	Andrea	Jones
Design Team	Bharath	Gupta
Software Group	Steve	Wells
Software Group	Santosh	Kumar
Software Group	Alice	James
Software Group	Ann	Jacob